

ACIS-X: An Autonomous, Event-Driven Multi-Agent Framework for Real-Time Credit Intelligence and Self-Healing Risk Assessment

Nikhil Nagendra¹, Pallavi K V¹, Johan K George¹, Nitin Bharadwaj MVS¹,
and Prasanna B Acharya¹

Department of Computer Science and Engineering,
AMC Engineering College, 18th K.M. Bannerghatta Main Road,
Bengaluru 560083, Karnataka, India
nikhilnag98@gmail.com

Abstract. Real-time risk analysis of credit for enterprises using financial systems requires the continuous and efficient processing of multiple data streams, including transactional activity, external financial indicators, and litigation-related signals. Current credit monitoring approaches typically rely on periodic batch analysis or manual review by credit control teams, thereby preventing early intervention when there is a decline in credit. This paper presents ACIS-X, an Autonomous Credit Intelligence System built as a distributed multi-agent framework upon Apache Kafka event streaming. ACIS-X decomposes the credit intelligence life-cycle into twelve specialized, loosely coupled agents spanning behavioral metrics computation, external signal enrichment, multi-tier risk fusion, policy enforcement, collections routing, and autonomous self-healing. A novel three-tier signal fusion model combines transactional behavioral metrics with external financial indicators and litigation intelligence into a unified risk score. Furthermore, the system guarantees regulatory compliance via SHAP-based model explainability persistence. The self-healing orchestration layer monitors agent health via heartbeat telemetry and autonomously restarts failed components with zero event loss, leveraging Kafka offset-based recovery. Experimental evaluation using 162 automated tests covering unit, integration, self-healing proof, and performance scenarios demonstrates: pipeline P95 latency under 500 ms across 500 events, mean time to recovery of under 30 s across 50 fault-injection cycles, exact-once deduplication verified over 10,000 repeated events, and near-linear horizontal throughput scaling exceeding $2.5\times$ baseline at four replicas. Risk score monotonicity, null-safety under partial signal availability, and thread-safe concurrent rebuilds are formally validated.

Keywords: Multi-agent systems · Event-driven architecture · Credit risk intelligence · Apache Kafka · Self-healing systems · Real-time analytics

1 Introduction

Enterprise credit risk management requires continuous assessment of customer payment health, external financial standing, and litigation exposure. Traditional systems operate on daily or weekly batch cycles, introducing latencies that prevent timely intervention when risk deteriorates intraday. The proliferation of event streaming platforms such as Apache Kafka [3] and advances in multi-agent software engineering [11] create an opportunity to fundamentally re-architect credit intelligence as a real-time, resilient, autonomous pipeline.

This paper presents ACIS-X (Autonomous Credit Intelligence System, Extended), a production-grade system implemented in Python 3.10 with Apache Kafka 3.6 as its event backbone. The system addresses three principal limitations of existing approaches:

1. **Temporal latency:** Batch systems miss intraday risk signal evolution critical for proactive collections.
2. **Signal siloing:** Financial, behavioral, and litigation risk signals are assessed in isolation, preventing holistic risk fusion.
3. **Operational fragility:** Single-process pipelines lack autonomous recovery, creating outage risk during agent failures.

The principal contributions of this work are:

- A twelve-agent distributed architecture decomposed across five functional layers, all communicating exclusively through Kafka topics.
- A three-tier signal fusion model combining on-time payment ratios, accounts-receivable aging, external financial indicators (debt/equity, ROE, PE ratio), and litigation severity into a composite risk score using validated weighting schemes.
- An autonomous self-healing orchestration layer with heartbeat-driven failure detection and Kafka offset-based zero-loss agent restart.
- Empirical validation through 162 automated tests: P95 pipeline latency < 500 ms, MTTR < 30 s under fault-injection cycles, exact-once deduplication over 10,000 events, and $\geq 2.5\times$ horizontal throughput scaling.

2 Related Work

Traditional credit scoring methods such as Altman’s Z-score [1] and logistic regression models [10] rely on annual financial statements, rendering them structurally incapable of intraday risk tracking. Ensemble methods [5] and LSTM-based approaches [2] improve accuracy but remain batch-oriented with hourly or daily refresh cycles.

Streaming financial systems demonstrate Kafka’s suitability for high-throughput financial processing [3], with fraud detection systems achieving sub-second latency [6]. However, comprehensive credit health monitoring combining

behavioral, financial, and litigation signals in a unified streaming pipeline has not been demonstrated.

Multi-agent systems have been applied to financial market simulation [12] and compliance monitoring, with Kumar et al. [4] providing foundational agent-oriented engineering methodology for FinTech. No prior work demonstrates MAS applied to streaming credit intelligence with integrated self-healing at production scale.

Self-healing systems grounded in the MAPE-K loop [9] and chaos engineering principles [8] provide coarse-grained recovery via container restart policies. Fine-grained, domain-aware agent restart protocols preserving event processing state – as implemented in ACIS-X – have not been demonstrated in financial AI contexts.

The critical gap addressed by ACIS-X lies at the intersection of real-time event streaming, multi-agent signal fusion, and autonomous self-healing, none of which prior works combine in a unified, production-validated system.

3 Problem Statement

Let $\mathcal{C} = \{c_1, \dots, c_N\}$ be enterprise customers in a credit relationship. For each customer c_i , the system must maintain a continuously updated risk score:

$$\mathcal{R}(c_i, t) = \mathcal{F}(\mathcal{B}(c_i, t), \mathcal{E}(c_i, t), \mathcal{L}(c_i, t)) \quad (1)$$

where $\mathcal{B}$ captures behavioral signals from transactional history, $\mathcal{E}$ captures external financial indicators, and $\mathcal{L}$ captures litigation exposure. The fusion function $\mathcal{F} : \mathbb{R}^p \times \mathbb{R}^q \times \mathbb{R}^r \rightarrow [0, 1]$ maps heterogeneous signals to a scalar risk score. Key constraints are: (1) end-to-end latency $\tau_{e2e} \leq \delta_{\max} = 1,000$ ms; (2) valid risk estimates under partial signal availability; (3) zero event loss under agent failure; and (4) idempotent processing of duplicate events.

4 System Architecture

ACIS-X is organized across five functional layers communicating exclusively through Apache Kafka, as shown in Fig. 1.




Fig. 1. ACIS-X five-layer multi-agent architecture. All inter-layer communication is mediated by Apache Kafka topics. The System Layer operates orthogonally across all processing layers.

4.1 Agent Inventory

Table 1. ACIS-X Agent Inventory – Subscriptions and Publications

Agent	Layer	Subscribes To	Publishes To
ScenarioGeneratorAgent	1	<code>acis.commands</code>	customers, invoices, payments
CustomerStateAgent	1	invoices, payments	<code>acis.metrics</code>
ExternalDataAgent	2	<code>acis.metrics</code>	<code>acis.metrics</code> (enriched)
ExternalScrapingAgent	2	<code>acis.metrics</code>	<code>acis.metrics</code> (litigation)
AggregatorAgent	2	<code>acis.metrics</code>	<code>acis.risk</code>
PaymentPredictionAgent	3	<code>acis.metrics</code>	<code>acis.predictions</code>
RiskScoringAgent	3	predictions, customers, metrics, risk	<code>acis.risk</code>
OverdueDetectionAgent	4	<code>acis.time</code>	<code>acis.invoices</code>
CreditPolicyAgent	4	<code>acis.risk</code>	<code>acis.collections</code>
CollectionsAgent	4	<code>acis.collections</code>	<code>acis.collections</code>
DBAgent	5	ALL topics	SQLite (persist)
QueryAgent / MemoryAgent	5	(synchronous calls)	(in-memory cache)

Table 1 summarizes the twelve agents, their Kafka subscriptions, and publications, derived directly from the implemented codebase.

4.2 Kafka Partitioning and Consumer Groups

Messages are keyed by `customer_id` hash, preserving per-customer ordering. Each agent type operates within its own dedicated consumer group (e.g., `customer-state-group`, `external-data-group`, `risk-scoring-group`, `collections-group`), enabling independent horizontal scaling per agent. Three partitions per topic allow up to three replicas of each agent to share a consumer group for load-balanced processing.

4.3 Persistence Layer

The SQLite schema comprises ten tables: `customers`, `invoices`, `payments`, `customer_risk_profile`, `customer_metrics`, `external_financials` (90-day TTL cache), `external_litigation` (90-day TTL cache), `risk_explanations` (SHAP audit log), `collections_log`, and `event_log`. DBAgent uses `INSERT OR REPLACE` for idempotent upserts. MemoryAgent maintains an in-memory customer state derived from database queries, providing fast access to aggregated metrics with thread-safe state management via `threading.Lock`.

5 Methodology and Mathematical Model

5.1 Three-Tier Signal Fusion

Tier 1 – Behavioral Signals ($\mathcal{B}$): Computed by CustomerStateAgent via a single bulk invoice query per event, with per-customer thread locking to prevent concurrent rebuild race conditions. The on-time ratio and average delay are:

$$\begin{aligned} \text{OTRatio}(c_i) &= \frac{|\{I_k : \text{paid}(I_k) \leq \text{due}(I_k)\}|}{|\mathcal{I}(c_i)|}, \\ \text{AvgDelay}(c_i) &= \frac{1}{|\mathcal{P}|} \sum_k \max((\text{paid}_k - \text{due}_k), 0) \end{aligned} \tag{2}$$

Tier 2 – External Signals ($\mathcal{E}, \mathcal{L}$): ExternalDataAgent queries financial APIs for debt/equity ratio, ROE, PE ratio, and sales growth, cached for 90 days. ExternalScrapingAgent queries legal news aggregation APIs with a 365-day lookback window, classifying litigation status as active, stayed, or resolved. Both agents publish to `acis.metrics` with `external.data.enriched` and `external.litigation.updated` event types respectively.

Tier 3 – Signal Fusion (AggregatorAgent): Financial and litigation signals are fused as:

$$\text{CombinedRisk}(c_i) = 0.6 \cdot \mathcal{E}(c_i) + 0.4 \cdot \mathcal{L}(c_i), \quad \text{CombinedRisk} \in [0, 1] \quad (3)$$

AggregatorAgent applies event freshness guards (discarding stale events older than a cached entry) and change-detection deduplication: a new **risk.profile.updated** event is published only when either financial or litigation risk values change from the last published values.

5.2 Payment Risk Prediction

PaymentPredictionAgent computes a per-invoice risk score using an explicit two-stage model consisting of a machine learning base and policy overrides. The first stage (r_{ml}) is a Random Forest ensemble:

$$r_{ml} = P(\text{Default}|X) = \frac{1}{M} \sum_{m=1}^M h_m(X) \quad (4)$$

where $M = 50$ is the number of decision trees, and the feature vector X incorporates both continuous behavioral metrics and context variables (e.g., s_{delay} , credit utilization, on-time ratio, litigation risk, and financial health).

To ensure systemic reliability, ACIS-X enforces this two-stage pipeline. The ML layer detects subtle shifts in continuous behavioral patterns, while the policy layer (r_{policy}) enforces non-negotiable risk penalties for catastrophic events (such as severe litigation or sudden credit rating drops):

$$r_{\text{final}} = \max(0, \min(1, r_{ml} + r_{\text{policy}})) \quad (5)$$

To ensure regulatory compliance and algorithmic transparency, ACIS-X computes exact Shapley Additive Explanations (SHAP) using TreeExplainer [7] for every prediction. The model decomposes the final risk score into an expected baseline value $E[f(X)]$ and additive feature attributions ϕ_i :

$$f(X) = E[f(X)] + \sum_{i=1}^N \phi_i \quad (6)$$

For example, a high-risk prediction (0.80) might be transparently decomposed as: *Baseline* (0.45) + *Payment Delay* (+0.20) + *High Utilization* (+0.10) + *Active Litigation* (+0.05). These SHAP values are persisted to the **risk_explanations** table, enabling auditors to trace exactly why a specific invoice was flagged for collections.

5.3 Risk Refinement and Final Scoring

RiskScoringAgent refines the base prediction using five normalized behavioral factors and an external context boost:

$$r_{\text{adj}} = r_{\text{pred}} + 0.85 \cdot \Delta_{\text{behavioral}} + 0.30 \cdot \Delta_{\text{temporal}} \quad (7)$$

where $\Delta_{\text{behavioral}}$ is defined as:

$$\begin{aligned} \Delta_{\text{behavioral}} = & 0.40 \cdot a_{\text{overdue}} + 0.20 \cdot a_{\text{payment}} \\ & + 0.15 \cdot a_{\text{delay}} + 0.15 \cdot a_{\text{outstanding}} + 0.10 \cdot a_{\text{exposure}} \end{aligned} \quad (8)$$

with each factor normalized to $[0, 1]$. The external context boost blends aggregated, financial, and litigation risks (weights 0.40, 0.30, 0.20, 0.10) and applies up to a 20% multiplicative boost to $\Delta_{\text{behavioral}}$. Temporal trend signals (velocity, stability, volatility) contribute Δ_{temporal} , detected via a `get_risk_velocity` query. The final score is clamped to $[0, 1]$ and classified: high > 0.75 , medium > 0.40 , low otherwise.

5.4 Self-Healing Protocol

The SelfHealingAgent implements a MAPE-K control loop. Each agent publishes a heartbeat event every $T_H = 5$ s. Agent a_j is declared failed when:

$$t_{\text{now}} - t_{\text{heartbeat}}(a_j) > \theta_{\text{timeout}} = 60 \text{ s} \quad (9)$$

Upon detection, the SelfHealingAgent emits an `agent.restart.requested` event, reinitializing the failed agent's Kafka consumer from its last committed offset. This guarantees zero message loss across all agent failure scenarios.

6 Implementation

Technology stack: Python 3.10, confluent-kafka 2.3, SQLite 3 (stdlib), pytest 7.4, Docker Compose (Kafka 3.6 + ZooKeeper 3.8). The full system comprises approximately 12 agent modules totalling over 350 KB of production Python code.

Lazy producer initialization: Kafka producers are created on the first `publish()` call per agent rather than at startup. This reduces startup Kafka connections from 20 to 2 (90% reduction) and startup time from > 60 s to < 30 s.

In-memory state management: MemoryAgent maintains an event-driven, in-memory view of customer state derived from database queries, with thread-safe access via `threading.Lock` and $O(1)$ idempotency checks using a hybrid Set+Deque guard (bounded to 10,000 entries).

Idempotent deduplication: DBAgent uses `INSERT OR REPLACE` with primary-key constraints. AggregatorAgent tracks last-published risk values per customer

in an `OrderedDict` (bounded to 10,000 entries) and skips publication when values are unchanged. `CollectionsAgent` maintains a set-based guard preventing duplicate collection actions for the same (customer, invoice, action) triple.

Cache management: Financial and litigation data are cached in SQLite with 90-day TTL. `AggregatorAgent` enforces a 24-hour in-memory cache TTL with periodic cleanup every 60 seconds, bounded to 50,000 customers. `RiskScoringAgent` context cache enforces 24-hour TTL with 5-minute cleanup intervals, bounded to 10,000 customers.

Thread safety: `CustomerStateAgent` uses per-customer `threading.Lock()` instances for rebuild operations, preventing concurrent rebuild race conditions that caused cache drift in earlier versions.

7 Experimental Evaluation

7.1 Test Infrastructure

All 162 automated tests execute offline (no live Kafka broker required) via `confluent-kafka` mock injection and `unittest.mock` patches for external APIs. Tests run under Python 3.10.11, `pytest` 7.4.3, on Windows 11 (Intel i7-12700H, 32 GB DDR5). The suite spans unit, integration, and performance categories across two test directories: `tests/regular_tests/` and `tests/suite/`.

7.2 Test Results Summary

Table 2. Complete Test Suite Results (162 Tests, All Pass)

Test Module	Property Verified	Tests	Result
test_unit_architecture_fixes	Agent contracts, schema integrity	22	PASS
test_aggregator_null_safety	Null/partial signal handling	5	PASS
test_risk_score_monotonicity	Risk increases with worsening inputs	4	PASS
test_exactly_once_processing	Duplicate event idempotency	2	PASS
test_throughput_scaling	Horizontal scaling $\geq 2.5\times$	1	PASS
test_self_healing_concurrent	Concurrent self-healing	1	PASS
suite/test_self_healing_proof	Agent failure recovery simulation	5	PASS
test_unit_self_healing	MAPE-K protocol correctness	5	PASS
test_unit_self_healing_health_routing	Health routing, escalation	8	PASS
test_recovery_action_correctness	Recovery action correctness	11	PASS
test_recovery_storm_prevention	Restart storm prevention	1	PASS
test_canonical_consumer_group_scaling	Consumer group partitioning	4	PASS
test_enrichment_ablation	Signal ablation (Tier 1/2/3)	2	PASS
test_ml_payment_prediction	Random Forest & 2-stage policy	3	PASS
test_query_client_thread_isolation	Thread-safe query isolation	3	PASS
test_overpayment_clamping	Risk score clamping to [0,1]	2	PASS
test_unit_circuit_breaker	Circuit breaker behavior	4	PASS
test_unit_kafka_client	Kafka client operations	7	PASS
test_unit_placement	Partition/placement engine	3	PASS
test_unit_external_agent	External API integration	4	PASS
suite/test_performance	P95 latency, MTTR, deduplication	3	PASS
suite/test_unit_core	Core agent contracts	15	PASS
suite/test_integration_contracts	Cross-agent event contracts	6	PASS
suite/test_lifecycle_trace	Invoice lifecycle, latency breakdown	4	PASS
suite/test_intelligence_comparison	Baseline comparison, ablation	3	PASS
suite/test_dynamic_discovery	Agent dynamic registration	2	PASS
Other unit modules	Schema, datetime, event, pipeline, query	29	PASS
Total		162	162 PASS

7.3 Performance Metrics

Table 3. Measured Performance Results from Automated Test Suite

Metric	Measured Value	Test / Source
Pipeline P95 latency (invoice $\rightarrow$ risk.scored)	< 500 ms	TestPipelineP95Latency (500 events)
Self-healing MTTR (mean, 50 cycles)	< 30 s	TestSelfHealingMTTR
Deduplication under 10,000 repeated events	Exactly 1 action	TestCollectionsAgentDuplicate
Horizontal throughput scaling (peak)	$\geq 2.5\times$ at $N = 4$	TestThroughputScaling
Risk score monotonicity	Verified	test_risk_score_monotonicity
Null-signal partial availability	Verified	test_aggregator_null_safety
Thread-safe concurrent rebuild	Verified	test_query_client_thread_isolation
Startup time (lazy producer init)	< 30 s	Architecture fix (20 $\rightarrow$ 2 connections)
External data cache TTL	90 days	ExternalDataAgent implementation
In-memory dedup cache bound	50,000 entries	Aggregator Agent

7.4 Comparison with Baseline Approaches

To evaluate the predictive performance of the Random Forest signal fusion model, we benchmarked it against standard industry baselines (Logistic Regression and XGBoost) using the synthetic dataset generated by the ACIS-X simulation environment. The ACIS-X Enriched Random Forest (which includes external litigation and financial signals) achieved an F1-score of 0.89, outperforming the Logistic Regression baseline (0.76) and an un-enriched XGBoost model (0.82) that relied solely on transactional data. This numerically validates that the inclusion of multi-tier external signals significantly improves predictive recall for high-risk defaults without sacrificing precision.

Table 4 compares ACIS-X against representative approaches on key architectural dimensions.

Table 4. Architectural Comparison with Related Systems

Property	Batch LR [10]	Streaming Fraud [6]	MAS Finance [12]	ACIS-X
Real-time processing	×	✓	×	✓
Multi-signal fusion	×	×	×	✓
Self-healing recovery	×	×	×	✓
Zero event loss on failure	×	×	N/A	✓
Horizontal scaling	×	✓	×	✓
Litigation integration	×	×	×	✓
Offline test coverage	×	×	×	162 tests

8 Discussion

Signal fusion contribution: The enrichment ablation test (`test_enrichment_ablation`) confirms that removing external financial and litigation signals degrades risk score accuracy, validating the three-tier fusion design. The AggregatorAgent’s freshness guard prevents stale cached values from overwriting more recent risk scores, a subtle correctness property verified by the null-safety test suite.

Architectural trade-offs: Choreography eliminates central orchestrator failure but complicates tracing, suggesting a need for OpenTelemetry span propagation. SQLite suffices locally but should be replaced by PostgreSQL for production, seamlessly supported by the modular DBAgent.

Self-healing correctness: The recovery storm prevention test confirms that the SelfHealingAgent does not trigger cascading restarts when multiple agents fail simultaneously. The health routing tests validate correct escalation paths across all agent states (healthy, degraded, timeout, failed).

9 Conclusion and Future Work

ACIS-X demonstrates that production-grade, real-time credit intelligence is achievable through principled multi-agent decomposition over Apache Kafka event

streaming. The system’s twelve-agent architecture, three-tier signal fusion model, SHAP-based regulatory explainability persistence, and MAPE-K self-healing layer address the temporal latency, signal siloing, and operational fragility limitations of existing batch-oriented credit systems. Validation through 162 automated tests confirms pipeline P95 latency under 500 ms, self-healing MTTR under 30s across robust agent failure simulations, exact-once deduplication, and near-linear horizontal scaling.

Future directions include temporal sequence modeling via LSTMs, privacy-preserving federated learning, and graph-based supply chain risk modeling.

Acknowledgement

The authors thank the faculty and laboratory infrastructure of the Department of Computer Science and Engineering, AMC Engineering College, Bengaluru, for supporting this research.

References

1. Altman, E.I.: Financial ratios, discriminant analysis and the prediction of corporate bankruptcy. *J. Finance* **23**(4), 589–609 (1968)
2. Chen, T., et al.: Streaming sequence modeling for real-time credit default prediction. *IEEE Access* **13**, 12050–12065 (2025)
3. Kreps, J., Narkhede, N., Rao, J.: Kafka: A distributed messaging system for log processing. In: *Proc. NetDB* (2011)
4. Kumar, A., et al.: Engineering autonomous agents for fintech. *IEEE Softw.* **41**(1), 45–52 (2024)
5. Li, Y., et al.: Ensemble learning for corporate credit risk evaluation. *IEEE Trans. Neural Netw. Learn. Syst.* **35**(5), 2450–2462 (2024)
6. Liu, H., et al.: Real-time fraud detection in financial streams. *Decis. Support Syst.* **182**, 114050 (2025)
7. Lundberg, S.M., et al.: From local explanations to global understanding with explainable AI for trees. *Nat. Mach. Intell.* **2**(1), 56–67 (2020)
8. Patel, K., et al.: Chaos engineering for resilient streaming analytics. *IEEE Internet Comput.* **28**(2), 22–29 (2024)
9. Singh, R., et al.: Self-healing microservices via MAPE-K in financial systems. *IEEE Trans. Cloud Comput.* **13**(1), 2310–2325 (2025)
10. Wang, X., et al.: Explainable machine learning in credit risk assessment: A SHAP-based approach. *Expert Syst. Appl.* **238**, 121751 (2024)
11. Zhang, L., et al.: Event-driven multi-agent architectures for autonomous finance. *IEEE Trans. Serv. Comput.* **17**(2), 112–126 (2024)
12. Zhao, J., et al.: Multi-agent reinforcement learning for financial market simulation. *J. Financ. Data Sci.* **6**(1), 15–30 (2024)